Lingyun Yang (杨凌云)

Updated September 20, 2024

Email: lyangbk@cse.ust.hk
GitHub: mental2008
LinkedIn: stephenyang1999
Phone: (+86) 135-0284-6103
Office: BDI 101, UC, HKUST
Web: https://www.lingyunyang.com/

Research Interests

I have a broad interest in resource management for large-scale data centers / AI infrastructure. Specifically, my research focuses on: (a) improving *resource efficiency* for AI/GPU clusters; (b) building *efficient* and *low-cost* AI model serving systems.

Education

Hong Kong University of Science and Technology (HKUST)
Department of Computer Science and Engineering
Ph.D. in Computer Science and Engineering — 2020 – Present
◇ Advisor: Prof. Wei Wang — (*expected to graduate in Fall 2025*)

South China University of Technology (SCUT)
School of Computer Science and Engineering
B.Eng. in Computer Science and Technology — 2016 – 2020
◇ Studied at All-English Innovation Class (GPA: 3.82/4); National Scholarship.

Internship

Alibaba Group — Hangzhou, China
Research Intern, Cluster Management Group — Dec. 2020 – Present
◇ Mentor: Dr. Yinghao Yu

Resource Management for AI/GPU Clusters

Mitigate GPU Resource Fragmentation

◇ [**ATC 2023**] Formally quantified *statistical GPU resource fragments* and proposed the *fragmentation gradient descent* scheduling algorithm to reduce resource fragmentation. Our scheduling policy can significantly reduce *unallocated* GPUs by up to 49% compared to state-of-the-art policies. [code] [trace]

◇ Proposed a GPU-disaggregated DLRM serving system to eliminate *resource mismatch* and meet *elastic* demand. By leveraging RDMA network to *separately* compute the computation graph on GPU and CPU nodes, it reduced CPU fragments by 53% and GPU fragments by 27%. During seasonal traffic peaks (e.g., Double 11 Shopping Festival), it saved up to 90% of GPUs when loaning GPU servers from training clusters.

◇ Developed ParaSet, a *best-effort* workload on Kubernetes that dynamically adjusts the number of instances and resource requirements based on the real-time resource availability in the cluster. It aims to fill resource fragments in the cluster and is integrated into KubeDL for internal use.

Large-Scale GPU Sharing in Production

◇ Enabled *large-scale GPU sharing* in production clusters, with over 10k shared GPU containers running daily. Support the co-location of GPU tasks with different priorities (e.g., *latency-sensitive*, *best-effort*). Designed and implemented the *node-level* agent and the *cluster-level* controller. The agent periodically collects and reports resource usage metrics and dynamically allocates GPU resources to containers. The controller calculates potential resource overcommitment and provides scheduling guidance to the cluster scheduler.

Efficient and Low-cost AI Model Serving Systems

Efficient Text-to-Image Diffusion Model Serving with Add-on Modules

◇ Developed SwiftDiffusion, a system that efficiently generates high-quality images with stable diffusion models and add-on modules (i.e., ControlNets and LoRAs). Incorporated serveral novel designs, including ControlNet-as-a-Service, asynchronous LoRA loading, and kernel optimization. Achieved up to 7.8× in latency and 1.6× in throughput without compromising image quality.

Auto-Configuration for AI Serving Service

◇ [**SoCC 2021**] Developed Morphling, an open-source auto-configuration framework for AI serving on Kubernetes. Combined *meta-learning* and *bayesian optimization* to quickly find the *optimal* configuration. It was widely used in Alibaba for automated recommendation of container resource specifications. [code]

Microsoft Research Asia (MSRA) Beijing, China
Research Intern, Innovation Engineering Group (IEG) Jul. 2019 – Jun. 2020
◇ Conducted research on model robustness, face recognition, attention mechanisms, knowledge distillation, and neural architecture search; Star of Tomorrow Internship Award of Excellence.

Publications

* *denotes co-first authors, sort in alphabetical order*

Refereed Papers in Conference Proceedings

[C3] Qizhen Weng*, **Lingyun Yang***, Yinghao Yu, Wei Wang, Xiaochuan Tang, Guodong Yang, Liping Zhang, "Beware of Fragmentation: Scheduling GPU-Sharing Workloads with Fragmentation Gradient Descent," in the *Proceedings of USENIX Annual Technical Conference* (**ATC '23**), Boston, MA, USA, July 2023. ***(CCF-A**, acceptance rate: 65/353=18.4%)*

[C2] Yongkang Zhang, Yinghao Yu, Wei Wang, Qiukai Chen, Jie Wu, Zuowei Zhang, Jiang Zhong, Tianchen Ding, Qizhen Weng, **Lingyun Yang**, Cheng Wang, Jian He, Guodong Yang, and Liping Zhang, "Workload Management in Alibaba Clusters: The Good, the Bad, and the Ugly," in the *Proceedings of ACM Symposium on Cloud Computing* (**SoCC '22**), San Francisco, CA, USA, November 2022. ***(CCF-B**, acceptance rate: 38/155=24.5%)*

[C1] Luping Wang*, **Lingyun Yang***, Yinghao Yu, Wei Wang, Bo Li, Xianchao Sun, Jian He, and Liping Zhang, "Morphling: Fast, Near-Optimal Auto-Configuration for Cloud-Native Model Serving," in the *Proceedings of ACM Symposium on Cloud Computing* (**SoCC '21**), Seattle, WA, USA, November 2021. ***(CCF-B**, acceptance rate: 46/145=31.7%)*

In submission / Preprint

[I2] Suyi Li*, **Lingyun Yang***, Xiaoxiao Jiang, Hanfeng Lu, Zhipeng Di, Weiyi Lu, Jiawei Chen, Kan Liu, Yinghao Yu, Tao Lan, Guodong Yang, Lin Qu, Liping Zhang, Wei Wang, "SwiftDiffusion: Efficient Diffusion Model Serving with Add-on Modules," *arXiv preprint arXiv:2407.02031*, 2024.

[I1] **Lingyun Yang**, Yongchen Wang, Yinghao Yu, Qizhen Weng, Jianbo Dong, Kan Liu, Chi Zhang, Yanyi Zi, Hao Li, Zechao Zhang, Nan Wang, Yu Dong, Menglei Zheng, Lanlan Xi, Xiaowei Lu, Liang Ye, Guodong Yang, Binzhang Fu, Tao Lan, Liping Zhang, Lin Qu, Wei Wang, "GPU-Disaggregated Serving for Deep Learning Recommendation Models at Scale," *under review.*

Awards

◇ Postgraduate Scholarship — 2020 – Present, HKUST
◇ Star of Tomorrow Internship Award of Excellence — Jul. 2020, MSRA
◇ Merit Student & Excellent Student Cadre — Nov. 2019, SCUT
◇ National Scholarship — Oct. 2019, China
◇ Silver Medal, ICPC China Xian National Invitational Contest — May 2019
◇ First Prize, 17th Guangdong Collegiate Programming Contest — May 2019
◇ Silver Medal, 37Games Cup Programming Contest — Apr. 2019
◇ Gold Medal, SCUT ACM Programming Contest — Apr. 2019
◇ Bronze Medal, ACM-ICPC Asia Xuzhou Regional Contest — Oct. 2018
◇ Silver Medal, 1st Xiao Mi Collegiate Programming Contest — Sept. 2018
◇ Gold Medal, SCUT ACM Programming Contest — Apr. 2018
◇ The First Prize Scholarship — Nov. 2017, SCUT
◇ Bronze Medal, ACM-ICPC Asia Xian Regional Contest — Oct. 2017
◇ Gold Medal, 12th China Youth Robot Competition — Jul. 2012
◇ Champion, RoboCup Youth Robot World Cup, China Division — Mar. 2012

Academic

Artifact Evaluation Committee

Services

◇ SIGCOMM (2024), HPCA (2024), SOSP (2023), OSDI (2023), ATC (2023), MLSys (2023)

External Reviewer

◇ INFOCOM (2022–2025), ICDCS (2023), APSys (2021), MSN (2021), Qshine (2020)

Student Helper

◇ APNet (2023), ICMLC & ICWAPR (2018)

Teaching Activities

Hong Kong University of Science and Technology

Teaching Assistant, Department of Computer Science and Engineering

◇ CSIT6000O: Advanced Cloud Computing (Spring 2022, Spring 2023)

◇ COMP4651: Cloud Computing and Big Data Systems (Spring 2021, Fall 2021, Spring 2024)

◇ COMP3511: Operating Systems (Fall 2023)

Other Experience

ACM-ICPC Competition Group — SCUT

Group Member & Team Leader — 2016 – 2019

◇ Coach: Prof. Chuhua Xian

◇ Major domains: Dynamic Programming, Number Theory, Data Structure, etc.

Machine Learning & Cybernetics Research Group — SCUT

Undergraduate Research Assistant — 2017 – 2019

◇ Advisor: Prof. Patrick Chan

◇ Projects: Fundus Stitching, Tableware Recognition, and NN Visualization.

Tencent Innovation Club — SCUT, CSE

Vice Chairman — 2018 – 2019

◇ Led the *largest* student club in SCUT CSE, sponsored by Tencent.

ByteDance Summer Camp — Beijing, China

Camper, Algorithm track — Aug. 2019

◇ Mentor: Dr. Yibo Zhu

◇ Totally 150 participants selected from more than 6k candidates.

Skills

Programming Languages: Golang, C++, Python, Javascript

Toolkits: Kubernetes, Docker, Grafana, Git, LaTeX, SQL, MarkDown

Languages: English (fluent), Mandarin (Native speaker), Cantonese (Intermediate)

Misc

Play basketball & badminton & squash, workout at the gym, foodie.

My paper reading notes are available at https://paper.lingyunyang.com/.